

Results

Results I

This section reports the compiled plans for both worked example methods. Every number below is produced by the thin analysis script (`scripts/methods_analysis.py`), which calls `run_all_gates` and `compile_method` from `src/methods_dsl/` and writes `output/data/compiled_plans.json`, `output/reports/gate_report.json`, and `output/reports/trust_chain_report.json`. Running the script regenerates every artifact this section references.

Compiled-plan summary I

Table 1: Compiled-plan summary from `output/data/compiled_plans.json`, generated by `compile_method` for each of 2 worked example methods.

Method	Steps	Target	Plan hash (first 12 hex chars)
PBSPreparation	5	human	313b9b17de98
SensorCalibrationSweep	4	automated	d89cced19be6


Compiled-plan summary II

[@tbl:compiled_plans] shows PBSPreparation (a manual, HUMAN-target bench preparation) alongside SensorCalibrationSweep (a mixed AUTOMATED/HUMAN instrument-calibration procedure) — the second example exists specifically to demonstrate the controlled vocabulary generalizing beyond wet-lab protocols, as [@sec:introduction] claims.

Step-count comparison I

`[@fig:step_counts]` plots the step count for each compiled method.

Step-count comparison II



Validation gates I

Across both methods, the analysis script tallies 8 of 8 staged-gate evaluations passing (`run_all_gates` $\times$ 2 methods $\times$ 4 gates each). Every gate result is written to `output/reports/gate_report.json` — neither method in this manuscript is hand-picked to pass; both worked examples are constructed to satisfy the structural, semantic, plan, and target gates by design, since `compile_method` raises `MethodValidationError` and halts the pipeline on any gate failure.

Determinism I

Recompiling each example method twice and comparing `plan_hash` values yields: **determinism check = Yes**. This is a live re-compilation comparison performed by `src/manuscript_variables.py` at manuscript-build time, not a value asserted once and then transcribed — the same property [`@sec:methodology`] claims for `compile_method` is checked again here, independently, against the live build.

Provenance demonstration I

`scripts/methods_analysis.py` appends a 3-record demonstration hash-chain for one value (`calibration_offset`) through DECLARED → CALIBRATED → VERIFIED tiers and writes the result of `verify_chain` to `output/reports/trust_chain_report.json`: **chain verified = Yes.**

Validation I

The results were validated through the zero-mock tests/ suite:

- ▶ **Library tests** assert exact gate outcomes, scheduling order, and plan-hash determinism against real Method fixtures, including one fixture per gate-failure mode.
- ▶ **Script test** runs `run_methods_analysis()` against a temporary output root and confirms real worklist/CSV/Mermaid/JSON artifacts, reports, and a real PNG figure are written.
- ▶ **Manuscript-variable test** confirms every generated-variable name used in `manuscript/*.md` is emitted by `generate_variables`.

All tests pass with coverage exceeding the 90% project gate, with no mocks.

Discussion I

The results confirm the pipeline end to end: both worked examples pass every staged gate, compile deterministically, and produce a stable plan hash across repeated builds. The same `src/methods_ds1/` functions back the analysis script, the test suite, and this manuscript — which is the architectural point of the exemplar. Because every number here is produced by a tested function and regenerated on demand, the prose describes structure and provenance rather than transcribing values that would drift the moment the example methods changed.